

Developing Operators for Kubernetes

by Pawel Kopiczko & Ross Fairbanks



Agenda

1. Operators Overview
2. Custom Resource Definition
3. Operator Task
4. Exercise 1. Implement using REST API
5. Exercise 2. Implement using client-go

Introductions



Tools

- Minikube 0.23

```
$ minikube start --kubernetes-version 'v1.8.0'
```

- Go 1.9
- Git
- Git clone giantswarm/operator-workshop

```
$ git clone https://github.com/giantswarm/operator-workshop  
$GOPATH/src/github.com/giantswarm/operator-workshop
```

Operators Overview



What is an Operator

- An application specific controller that extends the Kubernetes API
- Manages a CRD (Custom Resource Definition) which replace TPRs (Third Party Resources)
- CRD stores state and the operator performs operational tasks
- Pattern first proposed by CoreOS in November 2016

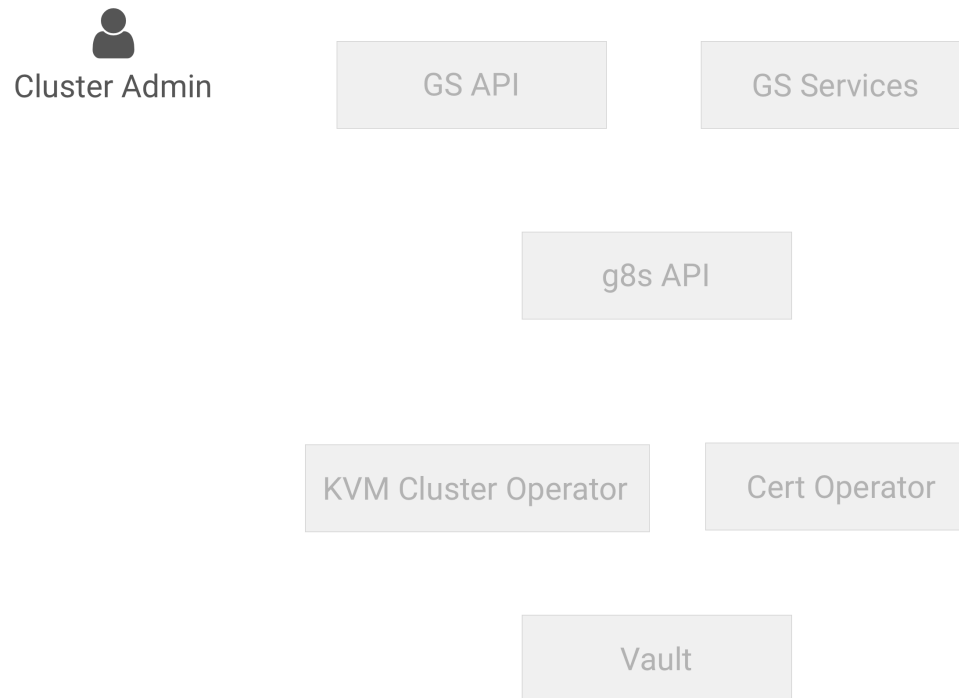
Example Operators

- Etcd
- Prometheus
- Rook

Giant Swarm Operators

- aws-operator
- kvm-operator
- cert-operator

Creating a guest cluster on g8s




 Kubernetes State

```
$ kubectl get pods
No resources found.
```

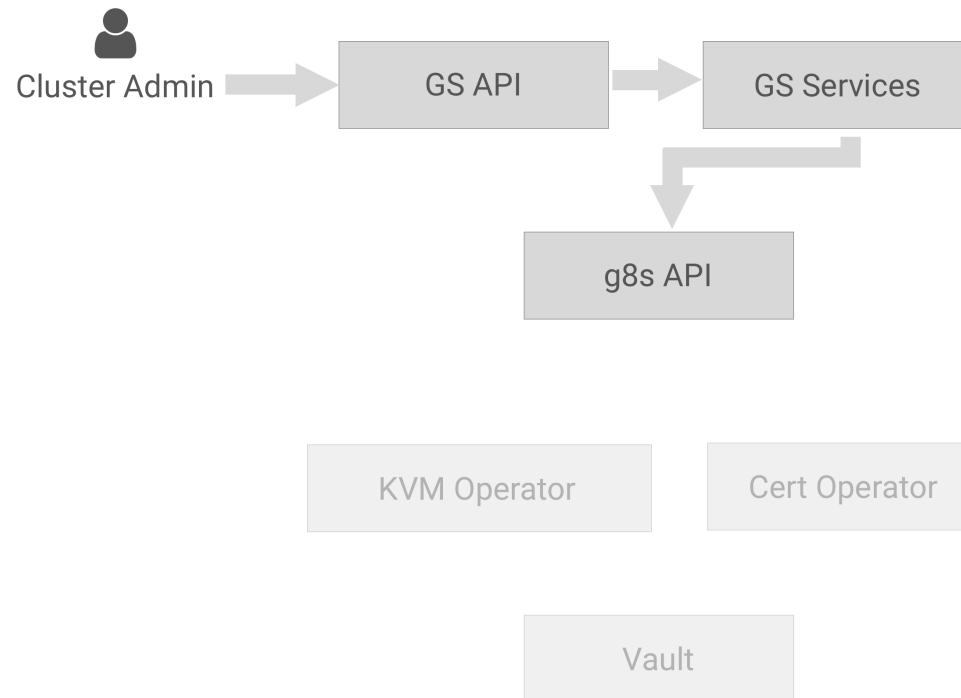
```
$ kubectl get namespaces
No resources found.
```

```
$ kubectl get secrets
No resources found.
```

 Custom Resource Definitions

No objects yet.

Cluster Admin performs POST to GS API




 Kubernetes State

```
$ kubectl get pods
No resources found.
```

```
$ kubectl get namespaces
```

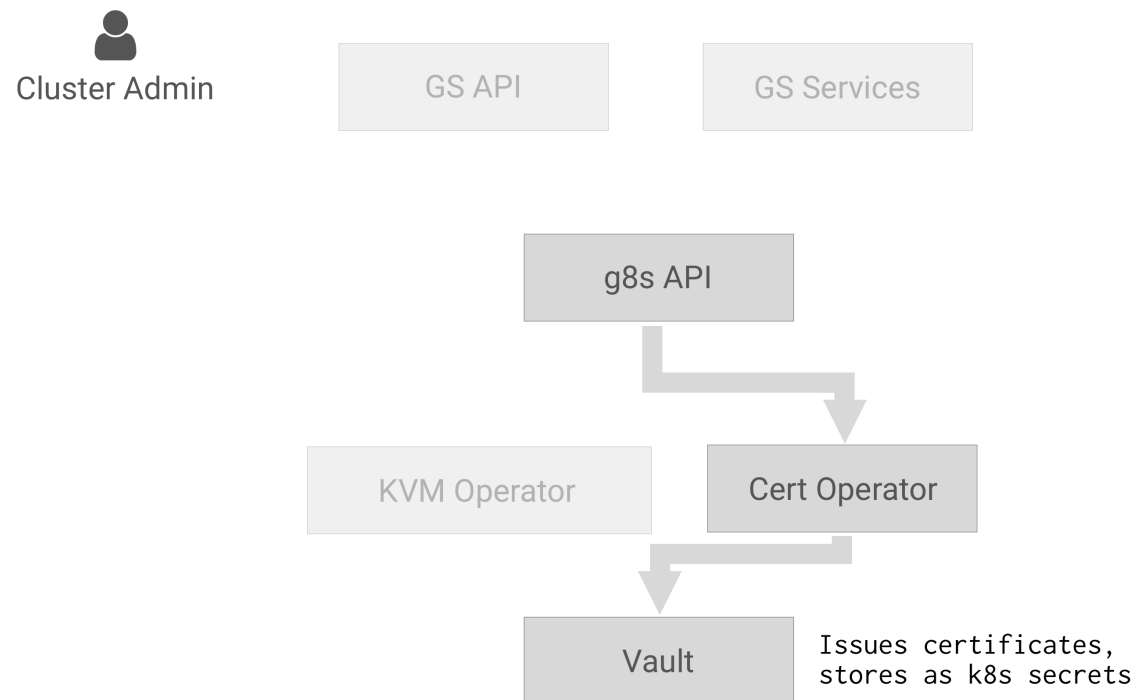
NAME	STATUS	AGE
uo91f	active	1m

```
$ kubectl get secrets
No resources found.
```

 Custom Resource Definitions

- Cluster object uo91f

Cert Operator acts on Cert CRDs



 Kubernetes State

```
$ kubectl get pods
No resources found.
```

```
$ kubectl get namespaces
```

NAME	STATUS	AGE
uo91f	active	1m

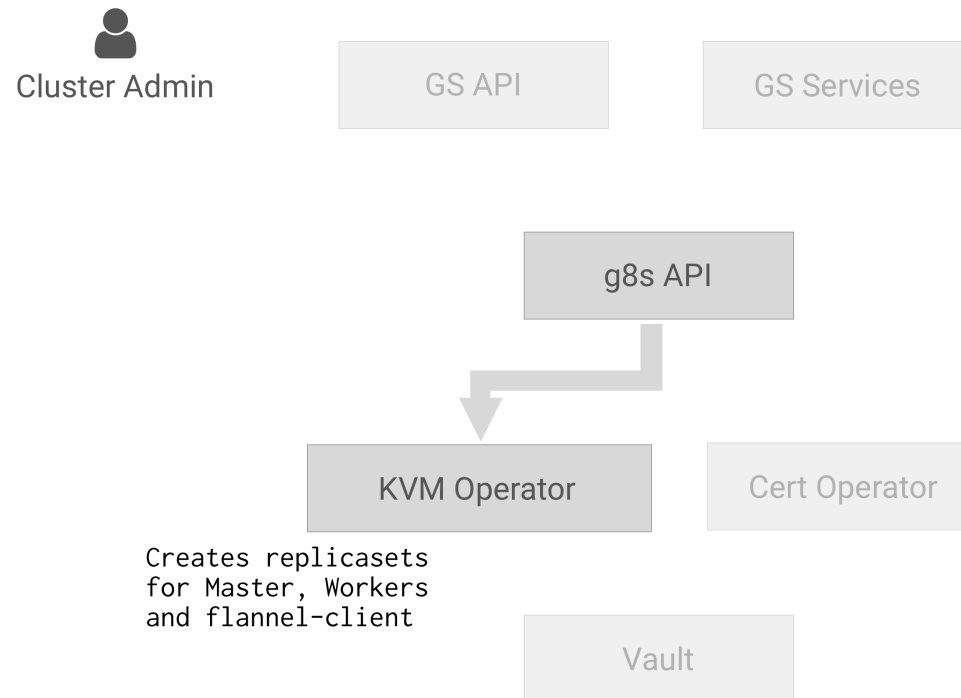
```
$ kubectl get secrets
```

NAME	
uo91f-root-ca	...
uo91f-cert-1	
...	

 Custom Resource Definitions

- Cluster object uo91f
- Multiple cert objects for components

Cluster Operator acts on Cluster CRDs



Kubernetes State

```
$ kubectl get pods
```

```
NAME
flannel-client-2256205...
flannel-client-2256205...
flannel-client-2256205...
master-3537008445-vnvxh
worker-2482331707-bx670
worker-2482331707-qq176
```

```
$ kubectl get namespaces
```

```
NAME    STATUS   AGE
uo91f   active   1m
```

```
$ kubectl get secrets
```

```
NAME      ...
uo91f-cert-1
...
```

Custom Resource Definition



Third Party Resources (Legacy)

- First mechanism for extending Kubernetes
- Added in Kubernetes 1.2
- Deprecated in Kubernetes 1.7
- Removed in Kubernetes 1.8

Custom Resource Definition

- Added in Kubernetes 1.7
- Reimplemented based on lessons learned from TPRs
- Namespace or cluster scoped
- Better pluralization support
- Shortname support
- Finalizers

Example CRD

```
apiVersion: apiextensions.k8s.io/v1beta1
kind: CustomResourceDefinition
metadata:
  name: crontabs.stable.example.com
spec:
  group: stable.example.com
  version: v1
  scope: Namespaced
  names:
    plural: crontabs
    singular: crontab
    kind: CronTab
    shortNames:
      - ct
```


Example Custom Resource

```
apiVersion: "stable.example.com/v1"
kind: CronTab
metadata:
  name: my-cron-object
spec:
  cronSpec: "* * * * */5"
  image: my-cron-image
```

API endpoints

- Core group endpoint (also known as legacy group)

```
/api/v1/pods
```

- Named group endpoint

```
/apis/apps/v1beta1/deployments
```

- CRD endpoint

```
/apis/stable.example.com/v1/crontabs
```

Demo



Finalizers

- Available in Kubernetes 1.8
- Allow controllers to implement asynchronous pre-delete hooks

```
apiVersion: "stable.example.com/v1"  
kind: CronTab  
metadata:  
  name: my-cron-object  
  finalizers:  
    - finalizer.stable.example.com
```

Validation

- Alpha in Kubernetes 1.8
- Allows validation of custom objects using OpenAPI V3 schema

```
spec:
  validation:
    # openAPIV3Schema is the schema for validating custom objects.
    openAPIV3Schema:
      properties:
        spec:
          properties:
            cronSpec:
            replicas:
              type: integer
              minimum: 1
              maximum: 10
```